

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
```

In [2]:

```
df = pd.read_csv("USA_Housing.csv")
df.head()
```

Out[2]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry 674\nLaurabury, 370
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Vi Suite 079\nL Kathleen, C
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Eliza Stravenue\nDaniel WI 064
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPC 44
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nF AE 09

In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                     5000 non-null   float64
1   Avg. Area House Age                  5000 non-null   float64
2   Avg. Area Number of Rooms            5000 non-null   float64
3   Avg. Area Number of Bedrooms         5000 non-null   float64
4   Area Population                      5000 non-null   float64
5   Price                               5000 non-null   float64
6   Address                             5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

In [4]:

```
df.describe()
```

Out[4]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
<b>count</b>	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
<b>mean</b>	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
<b>std</b>	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
<b>min</b>	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
<b>25%</b>	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
<b>50%</b>	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
<b>75%</b>	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
<b>max</b>	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06

In [5]:

```
df.corr().style.background_gradient() # to check the correlation without using a heatmap
```

Out[5]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
<b>Avg. Area Income</b>	1.000000	-0.002007	-0.011032	0.019788	-0.016234	0.639734
<b>Avg. Area House Age</b>	-0.002007	1.000000	-0.009428	0.006149	-0.018743	0.452543
<b>Avg. Area Number of Rooms</b>	-0.011032	-0.009428	1.000000	0.462695	0.002040	0.335664
<b>Avg. Area Number of Bedrooms</b>	0.019788	0.006149	0.462695	1.000000	-0.022168	0.171071
<b>Area Population</b>	-0.016234	-0.018743	0.002040	-0.022168	1.000000	0.408556
<b>Price</b>	0.639734	0.452543	0.335664	0.171071	0.408556	1.000000

In [7]:

```
df.corr()["Price"].sort_values()
```

Out[7]:

```
Avg. Area Number of Bedrooms    0.171071
Avg. Area Number of Rooms       0.335664
Area Population                  0.408556
Avg. Area House Age             0.452543
Avg. Area Income                0.639734
Price                           1.000000
Name: Price, dtype: float64
```

- By using this method we can find the correlation of prices as compared to the Average no. of bedrooms, Area population etc.

-Now by using iloc method we will divide the dataset into X(independent values) and Y(Dependent Values)

## Dividing Dataset into x and y

In [9]:

```
x = df.iloc[:, :-2].values
y = df.iloc[:, -2].values
# using the values method we can convert the data into numpy array
```

In [10]:

```
x
```

Out[10]:

```
array([[7.95454586e+04, 5.68286132e+00, 7.00918814e+00, 4.09000000e+00,
        2.30868005e+04],
       [7.92486425e+04, 6.00289981e+00, 6.73082102e+00, 3.09000000e+00,
        4.01730722e+04],
       [6.12870672e+04, 5.86588984e+00, 8.51272743e+00, 5.13000000e+00,
        3.68821594e+04],
       ...,
       [6.33906869e+04, 7.25059061e+00, 4.80508098e+00, 2.13000000e+00,
        3.32661455e+04],
       [6.80013312e+04, 5.53438842e+00, 7.13014386e+00, 5.44000000e+00,
        4.26256202e+04],
       [6.55105818e+04, 5.99230531e+00, 6.79233610e+00, 4.07000000e+00,
        4.65012838e+04]])
```

In [11]:

```
y
```

Out[11]:

```
array([1059033.55787012, 1505890.91484695, 1058987.98787608, ...,
       1030729.58315229, 1198656.87240769, 1298950.48026696])
```

## Creating the model (dividing training and testing set)

In [12]:

```
from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(x,y, test_size= 0.2, random_state=1)
```

## Performing Linear Regression by importing necessary Module

In [14]:

```
from sklearn.linear_model import LinearRegression
linreg = LinearRegression()
linreg.fit(xtrain,ytrain)
ypred = linreg.predict(xtest)
```

## Displaying Intercept

In [17]:

```
linreg.intercept_
```

Out[17]:

```
-2637185.6400765334
```

## Displaying Coefficient

In [18]:

```
linreg.coef_
```

Out[18]:

```
array([2.16667346e+01, 1.64990052e+05, 1.20784238e+05, 1.54252468e+03,
       1.51503697e+01])
```

**To interpret the Coefficient more easily we need to convert this data into a Dataframe**

In [20]:

```
x = df.iloc[:, :-2]
y = df.iloc[:, -2]
```

In [21]:

```
coef_df = pd.DataFrame(linreg.coef_,x.columns, columns = ["Coefficient"])
coef_df
```

Out[21]:

	Coefficient
Avg. Area Income	21.666735
Avg. Area House Age	164990.051829
Avg. Area Number of Rooms	120784.238317
Avg. Area Number of Bedrooms	1542.524676
Area Population	15.150370

## Now, interpreting the coefficient has become much more easier!!

- for every unit growth in Average Area income, we estimate that the price will increase upto \$21.66
- for every unit growth in Average area House age, we estimate that the price will increase upto \$164990.05
- for every unit growth in Average area no. of rooms, we estimate that the price will increase upto \$120784.23
- for every unit growth in Average area no. of bedrooms, we estimate that the price will increase upto \$1542.52
- for every unit growth in Area population, we estimate that the price will increase upto \$15.150


## Visualising the Data after training and testing

In [23]:

```
plt.scatter(ytest,ypred)
```

Out[23]:

<matplotlib.collections.PathCollection at 0x971b927190>




In [26]:

```
sns.distplot((ytest-ypred), bins = 50)
```

Out[26]:

<AxesSubplot:ylabel='Density'>



## Displaying evaluation Metrics

- It is used to evaluate the the performance and the accuracy of the model
- we need to import some modules for evaluation

In [28]:

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

- MAE or Mean Absolute error represents the average of the absolute difference between the actual values and the predicted values in the given Dataset

In [34]:

```
print(f"MeanAbsoluteError :- {mean_absolute_error(ytest,ypred)}")
```

MeanAbsoluteError :- 82494.73770124864

- Mean Squared Error or MSE represents the average of the squared difeference between the actual and the predicted values in the given Dataset

In [35]:

```
print(f"MeanSquaredError :- {mean_squared_error(ytest,ypred)}")
```

MeanSquaredError :- 10543597313.62453

- Root Mean Squared Error or RMSE represents the squareroot Mean Squared Error. It measures the Standard Deviation of the residuals

In [36]:

```
print(f"RootMeanSquaredError :- {np.sqrt(mean_squared_error(ytest,ypred))}")
```

RootMeanSquaredError :- 102682.02040096665

- r2\_score helps to determine the accuracy of the model

In [37]:

```
print(f"rSquared :- {(r2_score(ytest,ypred))}")
```

rSquared :- 0.9215935236936377

## CONCLUSION

- After evaluating this Dataset we can say that this model is 92% accurate and can be used for future predictions
- All the evaluation metrics has been displayed in the above Dataset

In [ ]: